

Part 2 Math grades for both schools

```
In [2]: import pandas as pd                #Importing the required package
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.model_selection import GridSearchCV
from timeit import default_timer as timer
from sklearn.linear_model import LinearRegression
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=DeprecationWarning)
warnings.simplefilter(action='ignore', category=UserWarning)

In [ ]: data = pd.read_csv('student-mat.csv', sep=';')
data.head()

In [ ]: data = data.drop(['G1', 'G2'], axis=1) #dropping columns G1 and G2 in order to get a m
data.head()

In [ ]: print("Number of rows:", len(data)) #Looking for number of rows in total

In [ ]: data.isna().sum() #Checking for Null Values

In [ ]: data.info()                #checking for Data types of each variables

In [ ]: for var in ['traveltime', 'studytime', 'Medu', 'Fedu', 'famrel', 'freetime', 'goout', 'Dalc', 'Palc', 'absences', 'G1', 'G2']:
    data[var] = data[var].astype('category')
data.info()

In [ ]: num_col = data._get_numeric_data().columns                #Getting the numerical dataset fr
cat_col = list(set(data.columns)-set(num_col))                #By removing the numerical datase
num_col = list(num_col)[0:3]

In [ ]: nominal_col = ['higher', 'Fjob', 'address', 'guardian', 'school', 'paid', 'sex', 'Mjob', 'ac
ordinal_col = np.setdiff1d(cat_col, nominal_col)

In [ ]: for col in ordinal_col:
    data[col] = data[col].astype(int)                # Changing Ordinal data to int

In [ ]: for col in data[cat_col]:
    print(col, data[col].unique())                # Checking for erroneous data in
```

```
In [ ]: data[num_col].describe()
```

```
In [ ]: for i in data[num_col]:  
        print(i,data[i].unique())
```

```
In [ ]: data['G3'].describe()
```

```
In [ ]: data['G3'].unique()
```

Exploratory Data Analytics

```
In [ ]: sns.pairplot(data, hue = 'school')
```

```
In [ ]: corrs = data.corr()  
fig, ax = plt.subplots(figsize=(15,10))  
sns.heatmap(corrs, annot=True, fmt='.2f', ax=ax)  
plt.show()
```

Based on observation, it is shown the following:

Fedu and Medu has moderate correlation

Walc and Dalc has moderate correlation

Walc and goout has weak correlation

```
In [ ]: sns.set(rc={'figure.figsize':(15,5)})  
sns.swarmplot(data=data, x='school', y='age')
```

Based on observation, it is shown that:

Students in GP are relatively younger than MS

```
In [ ]: ax = sns.boxplot(data=data, x='school', y='absences')  
ax.set_ylim([0, 35])
```

Based on observation, it is shown that:

GP has more absentees compared to MS

```
In [ ]: data.groupby(by='school').mean()
```

```
In [ ]: sns.boxplot(data=data, x="school", y="G3",  
                  palette=["m", "g"]  
                  )  
sns.set(rc={'figure.figsize':(8,5)})
```

Based on Observation on the boxplot, it is shown that:

In terms of median, 75th and 25th percentile, students in GP are relatively doing better than students in MS

Reducing Skewness in numerical data

```
In [ ]: skew_limit = 0.75

skew_cols = (data[num_col].skew()
             .sort_values(ascending=False)
             .to_frame()
             .rename(columns={0: 'Skew'})
             .query('abs(Skew) > {}'.format(skew_limit)))

skew_cols
```

```
In [ ]: field = "absences"
fig, (ax_original, ax_log1p) = plt.subplots(1, 2, figsize=(15, 5))

data[field].hist(ax=ax_original)

# Apply a log transformation (numpy syntax) to this column
data[field].apply(np.log1p).hist(ax=ax_log1p)

# Formatting of titles etc. for each subplot
ax_original.set(title='before np.log1p', ylabel='frequency', xlabel='value')
ax_log1p.set(title='after np.log1p', ylabel='frequency', xlabel='value')
fig.suptitle('Field "{}".format(field));
print('pop_ornigial skewness: ',data[field].skew())
print('pop_log1p skewness: ',data[field].apply(np.log1p).skew())
# fall in skewness after log1p
```

```
In [ ]: # apply log1p across all numerical columns

for col in data[skew_cols.index]:
    # drop famrel as skewness increase after log1p
    data[col] = data[col].apply(np.log1p)
new_skew_cols = (data[skew_cols.index].skew()
                .sort_values(ascending=False)
                .to_frame()
                .rename(columns={0: 'Skew'}))

skew_cols['New_skew'] = new_skew_cols
skew_cols['Difference'] = abs(skew_cols['New_skew']) - abs(skew_cols['Skew'])
skew_cols
```

MinMax scale numerical data

```
In [ ]: MinM = MinMaxScaler()
for col in num_col:
    data[col] = MinM.fit_transform(data[[col]])
```

One hot encoding of categorical Data

```
In [ ]: data = pd.get_dummies(data, columns=nominal_col,drop_first=True)
data.head()
# df is the transformed dataest for ML, df is the original dataset
```

Polynomial Features

```
In [ ]: from sklearn.preprocessing import PolynomialFeatures
```

```
In [ ]: feature_cols = list(filter(lambda x: x!= 'G3', data.columns))

polyf = PolynomialFeatures(degree=2, include_bias=False,)
X_pf = polyf.fit_transform(data[feature_cols])
```

Train Test Split

```
In [ ]: from sklearn.model_selection import train_test_split

train,test = train_test_split(data,
                              train_size = 0.7,
                              test_size = 0.3,
                              random_state = 42)

feature_cols = list(filter(lambda x: x!= 'G3', data.columns))

X_train = train[feature_cols]
y_train = train['G3']

X_test = test[feature_cols]
y_test = test['G3']
```

```
In [ ]: # splitting of Polynomial features
X_pf_train = X_pf[X_train.index]
y_pf_train = data['G3'][X_train.index]

X_pf_test = X_pf[X_test.index]
y_pf_test = data['G3'][X_test.index]
```

Model Kfold + Evaluation

```
In [ ]: # define root mean sq error func
def rmse(ytrue, ypredicted):
    return np.sqrt(mean_squared_error(ytrue, ypredicted))

folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
```

```
In [ ]: import statsmodels.api as sm
def diagnostic_plot(y_test, prediction):
    residual = y_test - prediction

    fig, axs = plt.subplots(2, 2, figsize=(15,10))
    fig.suptitle('Diagnostic plots')
    axs[0,0].plot([0, 15], [0, 20], ls="--", c="red", alpha=0.5)
    axs[0,0].scatter(y_test, prediction)
    axs[0,0].set_title('Truth-Prediction plot')
    axs[0,0].set_xlabel("Truth")
    axs[0,0].set_ylabel("Predictions")
```

```

sm.qqplot(residual,fit=True,line= 's', ax=axes[0,1])
axes[0,1].set_title('QQ plot')
axes[0,1].set_xlabel("Theoretical Quantiles")
axes[0,1].set_ylabel("Sample Quantiles")

sns.residplot(prediction,y_test,
               lowess=True,
               scatter_kws={'alpha': 0.5},
               line_kws={'color': 'red', 'lw': 1, 'alpha': 0.8},
               ax=axes[1,0])
axes[1,0].set_title('Residual plot')
axes[1,0].set_xlabel("Predicted")
axes[1,0].set_ylabel("Residual")

sqrt_standardized_residual=np.sqrt(np.abs(residual))
sns.regplot(prediction, sqrt_standardized_residual,
            scatter=True,
            lowess=True,
            line_kws={'color': 'red', 'lw': 1, 'alpha': 0.5},
            ax=axes[1,1])
axes[1,1].set_title('Scale-Location plot')
axes[1,1].set_xlabel("Predicted")
axes[1,1].set_ylabel("Sqrt Standarized residuals")

```

Dummy Model

```

In [ ]: average_G3 = np.mean(y_test)
        average_G3

```

```

In [ ]: start = timer()

        dummy_prediction =[]
        for row in range(len(y_test)):
            dummy_prediction.append(average_G3)

        print("R2 score: ",r2_score(y_test, dummy_prediction))
        print('RMSE: ', rmse(y_test,dummy_prediction))

        end = timer()
        dummy_time = end - start
        print('Time Elapsed (sec):', dummy_time)

```

Base Linear Regression

```

In [ ]: linear = LinearRegression()
        linear_cv = GridSearchCV(estimator = linear,
                                param_grid = {},
                                cv = folds)

        start = timer()
        linear_cv.fit(X_train, y_train)

        end = timer()
        linear_time = end - start
        print('Time Elapsed (sec):', linear_time)

```

```
In [ ]: linear_prediction = linear_cv.predict(X_test)

print("R2 score: ",r2_score(y_test, linear_prediction))
print('RMSE: ', rmse(y_test,linear_prediction))
```

```
In [ ]: diagnostic_plot(y_test, linear_prediction)
```

```
In [ ]: print("\u0332".join('Feature_coefficient (Top)'),'\n')
features = []
coeffs = []
coeffs_abs = []
for feature, coeff in zip(X_test.columns, linear_cv.best_estimator_.coef_):
    features.append(feature)
    coeffs.append(coeff)
    coeffs_abs.append(abs(coeff))

pd.DataFrame({'Features': features, 'Coefficients': coeffs , 'abs_coeffs':coeffs_abs})
```

```
In [ ]: print("\u0332".join('Feature_coefficient (Bottom)'),'\n')
pd.DataFrame({'Features': features, 'Coefficients': coeffs , 'abs_coeffs':coeffs_abs})
```

```
In [ ]: lm = LinearRegression()
lm_pf_cv = GridSearchCV(estimator = lm,
                        param_grid = {},
                        cv = folds)

start = timer()
lm_pf_cv.fit(X_pf_train, y_pf_train)

end = timer()
lm_pf_time = end - start
print('Time Elapsed (sec):', lm_pf_time)
```

```
In [ ]: lm_pf_prediction = lm_pf_cv.predict(X_pf_test)

print("R2 score: ",r2_score(y_pf_test, lm_pf_prediction))
print('RMSE: ', rmse(y_pf_test,lm_pf_prediction))
```

```
In [ ]: diagnostic_plot(y_test, lm_pf_prediction)
```

Lasso Regression

```
In [ ]: alphas = np.geomspace(0.001, 10, 30)
```

```
In [ ]: from sklearn.linear_model import LassoCV

start = timer()
lasso_estimator = LassoCV(alphas=alphas,
                          cv=5).fit(X_train,y_train)
lasso_prediction = lasso_estimator.predict(X_test)

print('Alpha param: ', lasso_estimator.alpha_)
print("R2 score: ",r2_score(y_test, lasso_prediction))
```

```
print('RMSE: ', rmse(y_test,lasso_prediction))

end = timer()
lasso_time = end - start
print('Time Elapsed (sec):', lasso_time)
```

```
In [ ]: print("Number of non-zero coeff: ", sum(lasso_estimator.coef_ != 0))
print("Mean coeff: ",sum(abs(lasso_estimator.coef_))/sum(lasso_estimator.coef_ != 0))
```

```
In [ ]: print("\u0332".join('Feature_coefficient (Top)'),'\n')
features = []
coeffs = []
coeffs_abs = []
for feature, coeff in zip(X_test.columns, lasso_estimator.coef_):
    features.append(feature)
    coeffs.append(coeff)
    coeffs_abs.append(abs(coeff))

pd.DataFrame({'Features': features, 'Coefficients': coeffs , 'abs_coeffs':coeffs_abs})
```

```
In [ ]: diagnostic_plot(y_test, lasso_prediction)
```

Ridge Regression

```
In [ ]: alphas = np.geomspace(1, 100, 30)
```

```
In [ ]: from sklearn.linear_model import RidgeCV

start = timer()
ridge_estimator = RidgeCV(alphas=alphas,
                           cv=5).fit(X_train,y_train)
ridge_prediction = ridge_estimator.predict(X_test)

print('Alpha param: ', ridge_estimator.alpha_)
print("R2 score: ",r2_score(y_test, ridge_prediction))
print('RMSE: ', rmse(y_test,ridge_prediction))

end = timer()
ridge_time = end - start
print('Time Elapsed (sec):', ridge_time)
```

```
In [ ]: print("Number of non-zero coeff: ", sum(ridge_estimator.coef_ != 0))
print("Mean coeff: ",sum(abs(ridge_estimator.coef_))/sum(ridge_estimator.coef_ != 0))
```

```
In [ ]: print("\u0332".join('Feature_coefficient (Top)'),'\n')
features = []
coeffs = []
coeffs_abs = []
for feature, coeff in zip(X_test.columns, ridge_estimator.coef_):
    features.append(feature)
    coeffs.append(coeff)
    coeffs_abs.append(abs(coeff))

pd.DataFrame({'Features': features, 'Coefficients': coeffs , 'abs_coeffs':coeffs_abs})
```

```
In [ ]: diagnostic_plot(y_test, ridge_prediction)
```

Elasticnet Regression

```
In [ ]: from sklearn.linear_model import ElasticNetCV
alphas = np.geomspace(0.01, 1, 30)
l1_ratios = np.linspace(0.1, 0.9, 9)

start = timer()
elasticNetCV = ElasticNetCV(alphas=alphas,
                           l1_ratio=l1_ratios,
                           cv=5).fit(X_train,y_train)
elasticNet_prediction = elasticNetCV.predict(X_test)

print('Alpha param: ', elasticNetCV.alpha_)
print('l1 ratio param: ', elasticNetCV.l1_ratio_)
print("R2 score: ",r2_score(y_test, elasticNet_prediction))
print('RMSE: ', rmse(y_test,elasticNet_prediction))

end = timer()
EN_time = end - start
print('Time Elapsed (sec):', EN_time)
```

```
In [ ]: print("Number of non-zero coeff: ", sum(elasticNetCV.coef_ != 0))
print("Mean coeff: ",sum(abs(elasticNetCV.coef_))/sum(elasticNetCV.coef_ != 0))
```

```
In [ ]: print("\u0332".join('Feature_coefficient (Top)'),'\n')
features = []
coeffs = []
coeffs_abs = []
for feature, coeff in zip(X_test.columns, elasticNetCV.coef_):
    features.append(feature)
    coeffs.append(coeff)
    coeffs_abs.append(abs(coeff))

pd.DataFrame({'Features': features, 'Coefficients': coeffs , 'abs_coeffs':coeffs_abs})
```

```
In [ ]: diagnostic_plot(y_test, elasticNet_prediction)
```

Stepwise regression - Forward

```
In [ ]: # Starting with no features, features are added iteratively based on the p-value
def forward_regression(X, y,
                      threshold_in,
                      verbose=False):
    initial_list = []
    included = list(initial_list)
    while True:
        changed=False
        exc = list(set(X.columns)-set(included))
        pval = pd.Series(index=exc)
        for column in exc:
            model = sm.OLS(y, sm.add_constant(pd.DataFrame(X[included+[column]]))).f
```



```

        pval[column] = model.pvalues[column]
    best = pval.min()
    if best < threshold_in:
        best_feature = pval.idxmin()
        included.append(best_feature)
        changed=True
    if verbose:
        print('Add {:17} with p-value {:.6}'.format(best_feature, best))

    if not changed:
        break

    return included

# Starting with all features, features are removed iteratively based on the p-value
def backward_regression(X, y,
                       threshold_out,
                       verbose=False):
    included=list(X.columns)
    while True:
        changed=False
        model = sm.OLS(y, sm.add_constant(pd.DataFrame(X[included]))).fit()
        # use all coefs except intercept
        pvalues = model.pvalues.iloc[1:]
        worst_pval = pvalues.max() # null if pvalues is empty
        if worst_pval > threshold_out:
            changed=True
            worst_feature = pvalues.idxmax()
            included.remove(worst_feature)
            if verbose:
                print('Drop {:17} with p-value {:.6}'.format(worst_feature, worst_pval))
        if not changed:
            break
    return included

```

```

In [ ]: start = timer()
forward_regression_features = forward_regression(X_train,y_train, threshold_in=0.05,
forward_regression_features

```

```

In [ ]: lm = LinearRegression()
forward = GridSearchCV(estimator = lm,
                       param_grid = {},
                       cv = folds)

# fit the model
forward.fit(X_train[forward_regression_features], y_train)

end = timer()
forward_time = end - start
print('Time Elapsed (sec):', forward_time)

```

```

In [ ]: forward_prediction = forward.predict(X_test[forward_regression_features])

print("R2 score: ",r2_score(y_test,forward_prediction))
print('RMSE: ', rmse(y_test,forward_prediction))

```

```

In [ ]: diagnostic_plot(y_test, forward_prediction)

```

Stepwise regression - Backward

```
In [ ]: start = timer()
backward_regression_features = backward_regression(X_train,y_train, threshold_out=0.
```

```
In [ ]: lm = LinearRegression()
backward = GridSearchCV(estimator = lm,
                        param_grid = {},
                        cv = folds)

# fit the model
backward.fit(X_train[backward_regression_features], y_train)

end = timer()
backward_time = end - start
print('Time Elapsed (sec):', backward_time)
```

```
In [ ]: backwards_prediction = backward.predict(X_test[backward_regression_features])

print("R2 score: ",r2_score(y_test, backwards_prediction))
print('RMSE: ', rmse(y_test,backwards_prediction))
```

```
In [ ]: diagnostic_plot(y_test, backwards_prediction)
```

Random Forest Regressor

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
random_tree = RandomForestRegressor(criterion = 'mse')
```

```
In [ ]: n_features = len(data.columns)
max_features = ['auto', 'sqrt', 'log2']
max_depth = [3, 4]
min_samples_split = [2, 5, 10]
min_samples_leaf = [1, 2, 4]
param_test = {'max_features': max_features,
              'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'min_samples_leaf': min_samples_leaf}
print(param_test)
```

```
In [ ]: random_tree = GridSearchCV(estimator = random_tree,
                                  param_grid = param_test,
                                  cv = folds)

start = timer()
random_tree.fit(X_train, y_train)
end = timer()
rf_time = end - start
print('Time Elapsed (sec):', rf_time)
```

```
In [ ]: random_tree_prediction = random_tree.predict(X_test)
print('Best param: ', random_tree.best_params_)
```

```
print("R2 score: ",r2_score(y_test, random_tree_prediction))  
print('RMSE: ', rmse(y_test,random_tree_prediction))
```

```
In [ ]: diagnostic_plot(y_test, random_tree_prediction)
```

```
In [ ]: from sklearn.tree import export_graphviz  
tree = random_tree.best_estimator_.estimators_[5]  
export_graphviz(tree, out_file='small.dot',  
                feature_names = X_train.columns)
```

```
In [ ]: from IPython.display import Image  
Image(filename = 'small.png')
```

```
In [ ]: feat_df = pd.DataFrame({'Feature':X_train.columns, 'Importance':random_tree.best_est  
feat_df = feat_df.sort_values(by = 'Importance', ascending=False)  
feat_df.head()
```

```
In [ ]: g = sns.barplot(data=feat_df, x='Feature', y= 'Importance')  
for item in g.get_xticklabels():  
    item.set_rotation(90)
```

```
In [ ]: models_pred = [dummy_prediction, linear_prediction, lm_pf_prediction, lasso_predicti  
time_elapsed = [dummy_time, linear_time, lm_pf_time, lasso_time, ridge_time, EN_time  
rmse_vals = []  
for pred in models_pred:  
    rmse_vals.append(rmse(y_test, pred))  
  
R2_score = []  
for pred in models_pred:  
    R2_score.append(r2_score(y_test, pred))  
  
labels = ['Dummy', 'Linear', 'Linear + PF', 'Ridge', 'Lasso', 'ElasticNet', 'Stepwise  
  
eval_df = pd.DataFrame({'RMSE':rmse_vals, 'R2 Score':R2_score, 'Time Elapsed':time_e  
eval_df
```

```
In [ ]:
```